



Reinforcement Learning Applied to the Shoals Marine Laboratory Smart Grid

Daniel Mattson

Advisor: Dr. Marek Petrik

Department of Computer Science, University of New Hampshire, Durham, NH 03824

Abstract

- The Shoals Marine Laboratory (SML) is a remote island lab with an isolated power grid.
- The island's goal is to solely use renewable energy, with a diesel generator as backup.
- Minimizing generator use and maximizing battery longevity is a complex task because of unpredictable environments and electricity demand.
- Reinforcement learning (RL) can be applied to this problem to improve generator usage beyond the naive policy the SML currently uses.
- Different underlying RL models are compared and evaluated in this project, with the linear spline approximation providing the best results.

Reinforcement Learning


- RL is a machine learning technique where an intelligent agent takes actions in an environment to maximize its notion of cumulative reward.
- The environment is modeled as a set of possible states and actions, and the agent receives a reward for each action it performs.
- The value function is used to quantify the perceived value of a state and is what the agent uses to choose actions. The goal is to be in states with high value.
- Approximating the value function is the primary problem that must be solved to choose an optimal action.

Data and Modeling

- All data used in this project was taken from the public SML website.
- Key variables to consider were the energy produced by solar, wind, and the generators on the island, along with the power usage from the battery storage.
- Fitted value iteration was used to approximate the value function using different features.
- Models using a neural net, linear spline, cubic spline, and radial basis functions were compared.

What the SML Currently Uses


Run the generator when under 70% charge



Value Function Approximations


Linear Spline

Fit using 1 week of data




Neural Network

Trained on 1 day of data



Radial Basis Function

Fit using 1 week of data



Results

- The value functions produced by all model types generally trend upward with higher charge, which is expected.
- The high values for lower charge seen in splines and RBFs could be explained by a lack of training data in that region.
- The neural network was only trained on 1 day of data due to long training times.
- With a larger training set or adjustments to the structure of the neural net better results are possible.
- The linear spline has the most correct approximation, matching the intuition of the problem.
- In the simulation, the linear spline approximation beats the naive policy currently used by the SML.

Conclusions

- The linear spline model performed the best out of all models in simulations.
- RL was effective when applied to optimizing SML smart grid operation.
- Total "cost" was significantly reduced.

Future work:

- Expand model to use predictions of future power demand and generation, time of year (season or week number) as inputs.
- Include a reward for having surplus energy in model.

References

SML Data: <https://sustainablesml.org/pages/systemList.php>

A. Ghasemi, A. Shojaeighadikolaei, K. Jones, M. Hashemi, A. G. Bardas and R. Ahmadi, "A Multi-Agent Deep Reinforcement Learning Approach for a Distributed Energy Marketplace in Smart Grids," 2020 IEEE International Conference on Communications, Control, and Computing Technologies for Smart Grids (SmartGridComm), 2020, pp. 1-6, doi: 10.1109/SmartGridComm47815.2020.9302981.

Y. Wan, J. Qin, X. Yu, T. Yang and Y. Kang, "Price-Based Residential Demand Response Management in Smart Grids: A Reinforcement Learning-Based Approach," in IEEE/CAA Journal of Automatica Sinica, vol. 9, no. 1, pp. 123-134, January 2022, doi: 10.1109/JAS.2021.1004287.